

# Use Azure OpenAI REST API

6 minutes

Azure OpenAI offers a REST API for interacting and generating responses that developers can use to add AI functionality to their applications. This unit covers example usage, input and output from the API.

## ⓘ Note

Before interacting with the API, you must create an Azure OpenAI resource in the Azure portal, deploy a model in that resource, and retrieve your endpoint and keys. Check out the [Getting started with Azure OpenAI Service](#) to learn how to do that.

For each call to the REST API, you need the endpoint and a key from your Azure OpenAI resource, and the name you gave for your deployed model. In the following examples, the following placeholders are used:

⌵ Expand table

Placeholder name	Value
YOUR_ENDPOINT_NAME	This base endpoint is found in the <b>Keys &amp; Endpoint</b> section in the Azure portal. It's the base endpoint of your resource, such as <code>https://sample.openai.azure.com/</code> .
YOUR_API_KEY	Keys are found in the <b>Keys &amp; Endpoint</b> section in the Azure portal. You can use either key for your resource.
YOUR_DEPLOYMENT_NAME	This deployment name is the name provided when you deployed your model in the Azure OpenAI Studio.

## Chat completions

Once you've deployed a model in your Azure OpenAI resource, you can send a prompt to the service using a POST request.

rest

```
curl https://YOUR_ENDPOINT_NAME.openai.azure.com/openai/deployments/YOUR_DE-  
PLOYMENT_NAME/chat/completions?api-version=2023-03-15-preview \  
-H "Content-Type: application/json" \  
-H "api-key: YOUR_API_KEY" \  
-d '{"messages":[{"role": "system", "content": "You are a helpful as-  
sistant, teaching people about AI."},  
{"role": "user", "content": "Does Azure OpenAI support multiple languages?"},  
{"role": "assistant", "content": "Yes, Azure OpenAI supports several lan-  
guages, and can translate between them."},  
{"role": "user", "content": "Do other Azure AI Services support translation  
too?"]}']'
```

The response from the API will be similar to the following JSON:

JSON

```
{  
  "id": "chatcmpl-6v7mkQj980V1yBec6ETrKPRqFjNw9",  
  "object": "chat.completion",  
  "created": 1679001781,  
  "model": "gpt-35-turbo",  
  "usage": {  
    "prompt_tokens": 95,  
    "completion_tokens": 84,  
    "total_tokens": 179  
  },  
  "choices": [  
    {  
      "message":  
        {  
          "role": "assistant",  
          "content": "Yes, other Azure AI Services also support  
translation. Azure AI Services offer translation between multiple languages  
for text, documents, or custom translation through Azure AI Services  
Translator."  
        },  
      "finish_reason": "stop",  
      "index": 0  
    }  
  ]  
}
```

REST endpoints allow for specifying other optional input parameters, such as `temperature`, `max_tokens` and more. If you'd like to include any of those parameters in your request, add them to the input data with the request.

## Embeddings

Embeddings are helpful for specific formats that are easily consumed by machine learning models. To generate embeddings from the input text, POST a request to the `embeddings` endpoint.

rest

```
curl https://YOUR_ENDPOINT_NAME.openai.azure.com/openai/deployments/YOUR_DE-  
PLOYMENT_NAME/embeddings?api-version=2022-12-01 \  
-H "Content-Type: application/json" \  
-H "api-key: YOUR_API_KEY" \  
-d '{"input": "The food was delicious and the waiter..."}'
```

When generating embeddings, be sure to use a model in Azure OpenAI meant for embeddings. Those models start with `text-embedding` or `text-similarity`, depending on what functionality you're looking for.

The response from the API will be similar to the following JSON:

JSON

```
{  
  "object": "list",  
  "data": [  
    {  
      "object": "embedding",  
      "embedding": [  
        0.0172990688066482523,  
        -0.0291879814639389515,  
        ....  
        0.0134544348834753042,  
      ],  
      "index": 0  
    }  
  ],  
  "model": "text-embedding-ada:002"  
}
```

## Next unit: Use Azure OpenAI SDK

Continue >

---